

# Deep Learning Project Part2

Group: 3

Members:

Juan I. Guimarey 5910690

Thiago Barbosa 5470137

- Images have been preprocessed on Part1
- We will load the data and start evaluating it
- We will define all our methods necessary
- Create the Model
- Train the Model
- Show predictions and accuracy results
- Show a Confusion Matrix of the results
- Show picture examples of incorrect and correct images
- Show a plot of Validations vs Losses

## Preprocessors

```
In [1]: #Preprocessors used  
from __future__ import absolute_import  
from __future__ import print_function  
from sklearn.metrics import confusion_matrix  
import matplotlib.pyplot as plt  
from datetime import timedelta  
from keras import regularizers  
import tensorflow as tf  
import seaborn as sns  
import numpy as np  
import h5py  
import time  
import os
```

Using TensorFlow backend.

## Import and load the data

```
In [2]: ▶ # default figure size for images
%matplotlib inline
plt.rcParams['figure.figsize'] = (16.0, 4.0)

h5f = h5py.File('SVHN_grey.h5', 'r')

# Load the training, test and validation set
X_train = h5f['X_train'][:]
y_train = h5f['y_train'][:]
X_test = h5f['X_test'][:]
y_test = h5f['y_test'][:]
X_val = h5f['X_val'][:]
y_val = h5f['y_val'][:]

h5f.close()

print('Training set', X_train.shape, y_train.shape)
print('Validation set', X_val.shape, y_val.shape)
print('Test set', X_test.shape, y_test.shape)
```

```
Training set (63733, 32, 32, 1) (63733, 10)
Validation set (9524, 32, 32, 1) (9524, 10)
Test set (26032, 32, 32, 1) (26032, 10)
```

## Defined Methods

```
In [3]: ▶ def get_batch(X, y, batch_size=512):
        for i in np.arange(0, y.shape[0], batch_size):
            end = min(X.shape[0], i + batch_size)
            yield(X[i:end], y[i:end])
```

```
In [4]: ▶ def plot_images(images, nrows, ncols, cls_true, mode=None):
        # Initialize the subplotgrid
        fig, axes = plt.subplots(nrows, ncols)

        # Randomly select nrows * ncols images
        rs = np.random.choice(images.shape[0], nrows*ncols)
        print("T = True, P = Prediction")
        for i, ax in zip(rs, axes.flat):

            if mode is None:
                title = "T: {0}".format(np.argmax(cls_true[i]))

            else:
                title = "T: {0}, P: {1}".format(np.argmax(cls_true[i]), mode[i])

            # Display the image
            ax.imshow(images[i, :, :, 0], cmap='binary')
            ax.set_title(title)
            ax.set_xticks([])
            ax.set_yticks([])
```

```
In [5]: ▶ def model(features):  
    weight_decay = 2e-4  
  
    # INPUT->[CONV->RELU->CONV->RELU->POOL]->DROPOUT->[FC->RELU]->FC  
    input_layer = tf.reshape(features, [-1, 32, 32, 1], name='Reshaped_Input')  
  
    conv1 = tf.layers.conv2d( inputs=input_layer, filters=32,  
        kernel_size=[5, 5], padding="same",  
        kernel_regularizer=regularizers.l2(weight_decay),  
        activation=tf.nn.relu)  
  
    pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2],  
        strides=2)  
  
    conv2 = tf.layers.conv2d( inputs=pool1, filters=64,  
        kernel_size=[5, 5], padding="same",  
        kernel_regularizer=regularizers.l2(weight_decay),  
        activation=tf.nn.relu)  
  
    pool2 = tf.layers.max_pooling2d(inputs=conv2, pool_size=[2, 2],  
        strides=2)  
  
    pool2_flat = tf.reshape(pool2, [-1, 8 * 8 * 64])  
    dense = tf.layers.dense(inputs=pool2_flat, units=256,  
        activation=tf.nn.relu)  
    dropout = tf.layers.dropout(inputs=dense, rate=0.8)  
  
    logits = tf.layers.dense(inputs=dropout, units=10)  
  
    return logits
```

## Building the Model

```
In [6]: ▶ max_epochs = 10  
    num_examples = X_train.shape[0]
```

```
In [7]: ▶ # Method to flatten results for Confusion Matrix  
    def flatten(lists):  
        flat = []  
        for numbers in lists:  
            for x in numbers:  
                flat.append(x)  
        return np.asarray(flat)
```

```
In [8]: tf.logging.set_verbosity(tf.logging.INFO)

# Our application logic will be added here
x = tf.placeholder(tf.float32, shape = [None, 32, 32, 1], name='Input_Data')
y = tf.placeholder(tf.float32, shape = [None, 10], name='Input_Labels')
y_cls = tf.argmax(y, 1)

discard_rate = tf.placeholder(tf.float32, name='Discard_rate')
os.environ['TF_CPP_MIN_LOG_LEVEL']='2'
```

```
In [9]: #with tf.name_scope('Model Prediction'):
prediction = model(x)
prediction_cls = tf.argmax(prediction, 1)
#with tf.name_scope('Loss'):
loss = tf.reduce_mean(tf.losses.softmax_cross_entropy(onehot_labels=y, logits=prediction))

optimizer = tf.train.AdamOptimizer().minimize(loss)
```

WARNING: Logging before flag parsing goes to stderr.  
W1203 17:09:33.333313 1888 deprecation.py:323] From <ipython-input-5-36aa2b28cf27>:9: conv2d (from tensorflow.python.layers.convolutional) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use `tf.keras.layers.Conv2D` instead.  
W1203 17:09:33.336306 1888 deprecation.py:506] From C:\Users\jguim\Anaconda3\lib\site-packages\tensorflow\python\ops\init\_ops.py:1251: calling VarianceScaling.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.  
Instructions for updating:  
Call initializer instance with the dtype argument instead of passing it to the constructor  
W1203 17:09:33.514797 1888 deprecation.py:323] From <ipython-input-5-36aa2b28cf27>:11: max\_pooling2d (from tensorflow.python.layers.pooling) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use keras.layers.MaxPooling2D instead.  
W1203 17:09:33.645449 1888 deprecation.py:323] From <ipython-input-5-36aa2b28cf27>:20: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use keras.layers.dense instead.  
W1203 17:09:33.954648 1888 deprecation.py:323] From <ipython-input-5-36aa2b28cf27>:21: dropout (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use keras.layers.dropout instead.  
W1203 17:09:34.063331 1888 deprecation.py:323] From C:\Users\jguim\Anaconda3\lib\site-packages\tensorflow\python\ops\losses\losses\_impl.py:121: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where

```
In [10]: ▶ # Predicted class equals the true class of each image?  
correct_prediction = tf.equal(prediction_cls, y_cls)  
  
# Cast predictions to float and calculate the mean  
accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
```

```
In [11]: ▶ sess = tf.Session()
```

```
In [12]: ▶ sess.run(tf.global_variables_initializer())
```

```
In [13]: ▶ batch_size = 512  
toss = 0.6  
accu = 0.0
```

```

In [14]: ▶ train_loss = []
valid_loss = []
start_time = time.time()
for epoch in range(max_epochs):
    print ('Training .....')
    epoch_loss = 0
    print ()
    print ('Epoch ', epoch+1 , ': ..... \n')
    step = 0

    ## Training epochs ....
    for (epoch_x , epoch_y) in get_batch(X_train, y_train, batch_size):
        _, train_accu, c = sess.run([optimizer, accuracy, loss],
            feed_dict={x: epoch_x, y: epoch_y, discard_rate: toss})
        train_loss.append(c)

        if(step%40 == 0):
            print ("Step:", step, ".....", "\nMini-Batch Loss   : ", c)
            print('Mini-Batch Accuracy :' , train_accu*100.0, '%')

            ## Validating prediction and summaries
            accu = 0.0
            for (epoch_x , epoch_y) in get_batch(X_val, y_val, 512):
                correct, _c = sess.run([correct_prediction, loss], feed_dict=
                    valid_loss.append(_c)
                accu+= np.sum(correct[correct == True])
            print('Validation Accuracy :' , accu*100.0/y_val.shape[0], '%')
            print ()
            step = step + 1

    print ('Epoch', epoch+1, 'completed out of ', max_epochs)

# Calculate net time
time_diff = time.time() - start_time

# Testing prediction and summaries

for (epoch_x , epoch_y) in get_batch(X_test, y_test, 512):
    correct = sess.run([correct_prediction],
        feed_dict={x: epoch_x, y: epoch_y, discard_rate: toss})
    accu+= np.sum(correct[correct == True])

print("Time usage: " + str(timedelta(seconds=int(round(time_diff)))))
print ("\n\n")

```

Step: 0 .....

Mini-Batch Loss : 0.25864834

Mini-Batch Accuracy : 92.1875 %

Validation Accuracy : 88.7757244855103 %

Step: 40 .....

Mini-Batch Loss : 0.13602126

Mini-Batch Accuracy : 96.2890625 %

Validation Accuracy : 88.62872742545149 %

```

Step: 80 .....
Mini-Batch Loss   : 0.14437835
Mini-Batch Accuracy : 94.53125 %
Validation Accuracy : 88.68122637547249 %

```

```


Step: 120 .....
Mini-Batch Loss   : 0.1806235
Mini-Batch Accuracy : 95.5078125 %
Validation Accuracy : 88.67072658546829 %

```

## Model Visualization

```
In [15]: plot_images(X_train, 3, 6, y_train);
```

T = True, P = Prediction



```
In [16]: test_pred = []
for (epoch_x , epoch_y) in get_batch(X_test, y_test, 512):
    correct = sess.run([prediction_cls],
        feed_dict={x: epoch_x, y: epoch_y, discard_rate: 0.0})
    test_pred.append((np.asarray(correct, dtype=int)).T)
```

## Confusion Matrix

```
In [17]: flat_array = flatten(test_pred)
flat_array = (flat_array.T)
flat_array = flat_array[0]
flat_array.shape
```


Out[17]: (26032,)

```
In [18]: #Set the figure size
plt.figure(figsize=(12, 8))

#Calculate the confusion matrix
cm = confusion_matrix(y_true=np.argmax(y_test, axis=1), y_pred=flat_array)

#Normalize the confusion matrix
cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis] * 100.0

#Show
sns.heatmap(cm, annot=True, cmap='Blues', fmt='.1f', square=True);
```



## Classified picture examples

First set of images are incorrect examples

Second set of images are correct examples




```
In [19]: ▶ # Find incorrectly classified examples
incorrect = flat_array != np.argmax(y_test, axis=1)

images = X_test[incorrect]
cls_true = y_test[incorrect]
mode = flat_array[incorrect]

plot_images(images, 3, 6, cls_true, mode);
```

T = True, P = Prediction




```
In [20]: ▶ # Find correctly classified examples
correct = np.invert(incorrect)

images = X_test[correct]
cls_true = y_test[correct]
mode = flat_array[correct]

plot_images(images, 3, 6, cls_true, mode);
```


T = True, P = Prediction



## Plot Validation vs Losses

```
In [21]: ▶ plt.style.use('classic')
plt.title("Validation and Training Losses")
plt.xlabel("Steps")
plt.ylabel("Values")
plt.plot(train_loss , 'r', label = 'loss')
plt.plot(valid_loss, 'b', label = 'validation')
plt.legend(loc='upper right', frameon=False)
```

Out[21]: <matplotlib.legend.Legend at 0x19348d375c0>



## Final Comments:

Project was achieved with an accuracy of 88.6% using 10 epochs A much higher accuracy can be achieved by adding more epochs but.. That would require higher computational power and time.

In [ ]: ▶

